

Serial Peripheral Interface (**SPI**)

SPI is the “Serial Peripheral Interface”, widely used with embedded systems because it is a simple and efficient interface: basically a multiplexed shift register. Its three signal wires hold a clock (SCK, often in the range of 1-20 MHz), a “Master Out, Slave In” (MOSI) data line, and a “Master In, Slave Out” (MISO) data line. SPI is a full duplex protocol; for each bit shifted out the MOSI line (one per clock) another is shifted in on the MISO line. Those bits are assembled into words of various sizes on the way to and from system memory. An additional chipselect line is usually active-low (nCS); four signals are normally used for each peripheral, plus sometimes an interrupt.

The **SPI** bus facilities listed here provide a generalized interface to declare SPI busses and devices, manage them according to the standard Linux driver model, and perform input/output operations. At this time, only “master” side interfaces are supported, where Linux talks to SPI peripherals and does not implement such a peripheral itself. (Interfaces to support implementing SPI slaves would necessarily look different.)

The programming interface is structured around two kinds of driver, and two kinds of device. A “Controller Driver” abstracts the controller hardware, which may be as simple as a set of GPIO pins or as complex as a pair of FIFOs connected to dual DMA engines on the other side of the **SPI** shift register (maximizing throughput). Such drivers bridge between whatever bus they sit on (often the platform bus) and **SPI**, and expose the **SPI** side of their device as a `struct spi_master`. **SPI** devices are children of that master, represented as a `struct spi_device` and manufactured from `struct spi_board_info` descriptors which are usually provided by board-specific initialization code. A `struct spi_driver` is called a “Protocol Driver”, and is bound to a `spi_device` using normal driver model calls.

The I/O model is a set of queued messages. Protocol drivers submit one or more `struct spi_message` objects, which are processed and completed asynchronously. (There are synchronous wrappers, however.) Messages are built from one or more `struct spi_transfer` objects, each of which wraps a full duplex **SPI** transfer. A variety of protocol tweaking options are needed, because different chips adopt very different policies for how they use the bits transferred with SPI.

```
struct spi_statistics
```

statistics for **spi** transfers

Definition

```

struct spi_statistics {
    spinlock_t lock;
    unsigned long    messages;
    unsigned long    transfers;
    unsigned long    errors;
    unsigned long    timedout;
    unsigned long    spi_sync;
    unsigned long    spi_sync_immediate;
    unsigned long    spi_async;
    unsigned long long bytes;
    unsigned long long bytes_rx;
    unsigned long long bytes_tx;
#define SPI_STATISTICS_HISTO_SIZE 17;
    unsigned long transfer_bytes_histo[SPI_STATISTICS_HISTO_SIZE];
    unsigned long transfers_split_maxsize;
};

```

Members

lock

lock protecting this structure

messages

number of **spi**-messages handled

transfers

number of **spi**_transfers handled

errors

number of errors during **spi**_transfer

timedout

number of timeouts during **spi**_transfer

spi_sync

number of times **spi**_sync is used

spi_sync_immediate

number of times **spi**_sync is executed immediately in calling context without queuing and scheduling

spi_async

number of times **spi**_async is used

bytes

number of bytes transferred to/from device

bytes_rx

number of bytes received from device

bytes_tx

number of bytes sent to device

transfer_bytes_histo

transfer bytes histogramm

transfers_split_maxsize

number of transfers that have been split because of maxsize limit

struct spi_device

Controller side proxy for an **SPI** slave device

Definition

```
struct spi_device {
    struct device      dev;
    struct spi_controller *controller;
    struct spi_controller *master;
    u32 max_speed_hz;
    u8 chip_select;
    u8 bits_per_word;
    u32 mode;
#define SPI_CPHA      0x01
#define SPI_CPOL      0x02
#define SPI_MODE_0    (0|0)
#define SPI_MODE_1    (0|SPI_CPHA);
#define SPI_MODE_2    (SPI_CPOL|0);
#define SPI_MODE_3    (SPI_CPOL|SPI_CPHA);
#define SPI_CS_HIGH   0x04
#define SPI_LSB_FIRST 0x08
#define SPI_3WIRE      0x10
#define SPI_LOOP       0x20
#define SPI_NO_CS      0x40
#define SPI_READY      0x80
#define SPI_TX_DUAL     0x100
#define SPI_TX_QUAD     0x200
#define SPI_RX_DUAL     0x400
#define SPI_RX_QUAD     0x800
#define SPI_CS_WORD     0x1000
#define SPI_TX_OCTAL    0x2000
#define SPI_RX_OCTAL    0x4000
#define SPI_3WIRE_HIZ   0x8000
    int irq;
    void *controller_state;
    void *controller_data;
    char modalias[SPI_NAME_SIZE];
    const char *driver_override;
    int cs_gpio;
    struct gpio_desc *cs_gpiod;
    uint8_t word_delay_usecs;
    struct spi_statistics statistics;
};
```

Members

dev

Driver model representation of the device.

controller

SPI controller used with the device.

master

Copy of controller, for backwards compatibility.

max_speed_hz

Maximum clock rate to be used with this chip (on this board); may be changed by the device's driver. The **spi**_transfer.speed_hz can override this for each transfer.

chip_select

Chipselect, distinguishing chips handled by **controller**.

bits_per_word

Data transfers involve one or more words; word sizes like eight or 12 bits are common. In-memory wordsizes are powers of two bytes (e.g. 20 bit samples use 32 bits). This may be changed by the device's driver, or left at the default (0) indicating protocol words are eight bit bytes. The **spi**_transfer.bits_per_word can override this for each transfer.

mode

The **spi** mode defines how data is clocked out and in. This may be changed by the device's driver. The "active low" default for chipselect mode can be overridden (by specifying SPI_CS_HIGH) as can the "MSB first" default for each word in a transfer (by specifying SPI_LSB_FIRST).

irq

Negative, or the number passed to `request_irq()` to receive interrupts from this device.

controller_state

Controller's runtime state

controller_data

Board-specific definitions for controller, such as FIFO initialization parameters; from `board_info.controller_data`

modalias

Name of the driver to use with this device, or an alias for that name. This appears in the sysfs "modalias" attribute for driver coldplugging, and in uevents used for hotplugging

cs_gpio

LEGACY: gpio number of the chipselect line (optional, -ENOENT when not using a GPIO line) use `cs_gpiod` in new drivers by opting in on the **spi** master.

cs_gpiod

gpio descriptor of the chipselect line (optional, NULL when not using a GPIO line)

word_delay_usecs

microsecond delay to be inserted between consecutive words of a transfer

statistics

statistics for the **spi** device

Description

A **spi_device** is used to interchange data between an **SPI** slave (usually a discrete chip) and CPU memory.

In **dev**, the `platform_data` is used to hold information about this device that's meaningful to the device's protocol driver, but not to its controller. One example might be an identifier for a chip variant with slightly different functionality; another might be information about how this particular board wires the chip's pins.

struct spi_driver

Host side "protocol" driver

Definition

```
struct spi_driver {
    const struct spi_device_id *id_table;
    int (*probe)(struct spi_device *spi);
    int (*remove)(struct spi_device *spi);
    void (*shutdown)(struct spi_device *spi);
    struct device_driver driver;
};
```

Members

`id_table`

List of **SPI** devices supported by this driver

`probe`

Binds this driver to the **spi** device. Drivers can verify that the device is actually present, and may need to configure characteristics (such as `bits_per_word`) which weren't needed for the initial configuration done during system setup.

`remove`

Unbinds this driver from the **spi** device

`shutdown`

Standard shutdown callback used during system state transitions such as powerdown/halt and kexec

`driver`

SPI device drivers should initialize the name and owner field of this structure.

Description

This represents the kind of device driver that uses **SPI** messages to interact with the hardware at the other end of a SPI link. It's called a "protocol" driver because it works through messages rather than talking directly to SPI hardware (which is what the underlying SPI controller driver does to pass those messages). These protocols are defined in the specification for the device(s) supported by the driver.

As a rule, those device protocols represent the lowest level interface supported by a driver, and it will support upper level interfaces too. Examples of such upper levels include frameworks like MTD, networking, MMC, RTC, filesystem character device nodes, and hardware monitoring.

```
void spi_unregister_driver(struct spi_driver *sdrv)
```

reverse effect of **spi**_register_driver

Parameters

```
struct spi_driver * sdrv
```

the driver to unregister

Context

can sleep

```
module_spi_driver(__spi_driver)
```

Helper macro for registering a **SPI** driver

Parameters

```
__spi_driver
```

spi_driver struct

Description

Helper macro for **SPI** drivers which do not do anything special in module init/exit. This eliminates a lot of boilerplate. Each module may only use this macro once, and calling it replaces

`module_init()` and `module_exit()`

```
struct spi_controller
```

interface to **SPI** master or slave controller

Definition

```

struct spi_controller {
    struct device dev;
    struct list_head list;
    s16 bus_num;
    u16 num_chipselect;
    u16 dma_alignment;
    u32 mode_bits;
    u32 bits_per_word_mask;
#define SPI_BPW_MASK(bits) BIT((bits) - 1);
#define SPI_BPW_RANGE_MASK(min, max) GENMASK((max) - 1, (min) - 1);
    u32 min_speed_hz;
    u32 max_speed_hz;
    u16 flags;
#define SPI_CONTROLLER_HALF_DUPLEX BIT(0) ;
#define SPI_CONTROLLER_NO_RX BIT(1) ;
#define SPI_CONTROLLER_NO_TX BIT(2) ;
#define SPI_CONTROLLER_MUST_RX BIT(3) ;
#define SPI_CONTROLLER_MUST_TX BIT(4) ;
#define SPI_MASTER_GPIO_SS BIT(5) ;
    bool slave;
    size_t (*max_transfer_size)(struct spi_device *spi);
    size_t (*max_message_size)(struct spi_device *spi);
    struct mutex io_mutex;
    spinlock_t bus_lock_spinlock;
    struct mutex bus_lock_mutex;
    bool bus_lock_flag;
    int (*setup)(struct spi_device *spi);
    void (*set_cs_timing)(struct spi_device *spi, u8 setup_clk_cycles, u8
hold_clk_cycles, u8 inactive_clk_cycles);
    int (*transfer)(struct spi_device *spi, struct spi_message *mesg);
    void (*cleanup)(struct spi_device *spi);
    bool (*can_dma)(struct spi_controller *ctlr, struct spi_device *spi, struct
spi_transfer *xfer);
    bool queued;
    struct kthread_worker kworker;
    struct task_struct *kworker_task;
    struct kthread_work pump_messages;
    spinlock_t queue_lock;
    struct list_head queue;
    struct spi_message *cur_msg;
    bool idling;
    bool busy;
    bool running;
    bool rt;
    bool auto_runtime_pm;
    bool cur_msg_prepared;
    bool cur_msg_mapped;
    struct completion xfer_completion;
    size_t max_dma_len;
    int (*prepare_transfer_hardware)(struct spi_controller *ctlr);
    int (*transfer_one_message)(struct spi_controller *ctlr, struct spi_message
*mesg);
    int (*unprepare_transfer_hardware)(struct spi_controller *ctlr);
    int (*prepare_message)(struct spi_controller *ctlr, struct spi_message
*message);
    int (*unprepare_message)(struct spi_controller *ctlr, struct spi_message
*message);
    int (*slave_abort)(struct spi_controller *ctlr);
    void (*set_cs)(struct spi_device *spi, bool enable);
    int (*transfer_one)(struct spi_controller *ctlr, struct spi_device *spi,
struct spi_transfer *transfer);
    void (*handle_err)(struct spi_controller *ctlr, struct spi_message *message);
    const struct spi_controller_mem_ops *mem_ops;
    int *cs_gpios;
    struct gpio_desc **cs_gpiods;
    bool use_gpio_descriptors;
    struct spi_statistics statistics;
    struct dma_chan *dma_tx;
    struct dma_chan *dma_rx;
    void *dummy_rx;
    void *dummy_tx;

```

```
int (*fw_translate_cs)(struct spi_controller *ctlr, unsigned cs);  
};
```

Members

dev

device interface to this driver

list

link with the global **spi_controller** list

bus_num

board-specific (and often SOC-specific) identifier for a given **SPI** controller.

num_chipselect

chipselects are used to distinguish individual **SPI** slaves, and are numbered from zero to num_chipselects. each slave has a chipselect signal, but it's common that not every chipselect is connected to a slave.

dma_alignment

SPI controller constraint on DMA buffers alignment.

mode_bits

flags understood by this controller driver

bits_per_word_mask

A mask indicating which values of bits_per_word are supported by the driver. Bit n indicates that a bits_per_word n+1 is supported. If set, the **SPI** core will reject any transfer with an unsupported bits_per_word. If not set, this value is simply ignored, and it's up to the individual driver to perform any validation.

min_speed_hz

Lowest supported transfer speed

max_speed_hz

Highest supported transfer speed

flags

other constraints relevant to this driver

slave

indicates that this is an **SPI** slave controller

max_transfer_size

function that returns the max transfer size for a **spi_device**; may be **NULL**, so the default **SIZE_MAX** will be used.

max_message_size

function that returns the max message size for a **spi_device**; may be **NULL**, so the default **SIZE_MAX** will be used.

io_mutex

mutex for physical bus access

bus_lock_spi_nlock

spi_nlock for SPI bus locking

bus_lock_mutex

mutex for exclusion of multiple callers

bus_lock_flag

indicates that the **SPI** bus is locked for exclusive use

setup

updates the device mode and clocking records used by a device's **SPI** controller; protocol code may call this. This must fail if an unrecognized or unsupported mode is requested. It's always safe to call this unless transfers are pending on the device whose settings are being modified.

set_cs_timing

optional hook for **SPI** devices to request SPI master controller for configuring specific CS setup time, hold time and inactive delay in terms of clock counts

transfer

adds a message to the controller's transfer queue.

cleanup

frees controller-specific state

can_dma

determine whether this controller supports DMA

queued

whether this controller is providing an internal message queue

kworker

thread struct for message pump

kworker_task

pointer to task for message pump kworker thread

pump_messages

work struct for scheduling work to the message pump

queue_lock

spi nlock to synchronise access to message queue

queue

message queue

cur_msg

the currently in-flight message

idling

the device is entering idle state

busy

message pump is busy

running

message pump is running

rt

whether this queue is set to run as a realtime task

auto_runtime_pm

the core should ensure a runtime PM reference is held while the hardware is prepared, using the parent device for the **spi** dev

cur_msg_prepared

spi _prepare_message was called for the currently in-flight message

cur_msg_mapped

message has been mapped for DMA

xfer_completion

used by core `transfer_one_message()`

max_dma_len

Maximum length of a DMA transfer for the device.

prepare_transfer_hardware

a message will soon arrive from the queue so the subsystem requests the driver to prepare the transfer hardware by issuing this call

transfer_one_message

the subsystem calls the driver to transfer a single message while queuing transfers that arrive in the meantime. When the driver is finished with this message, it must call

`spi_finalize_current_message()` so the subsystem can issue the next message

unprepare_transfer_hardware

there are currently no more messages on the queue so the subsystem notifies the driver that it may relax the hardware by issuing this call

prepare_message

set up the controller to transfer a single message, for example doing DMA mapping. Called from threaded context.

unprepare_message

undo any work done by `prepare_message()`.

slave_abort

abort the ongoing transfer request on an **SPI** slave controller

set_cs

set the logic level of the chip select line. May be called from interrupt context.

transfer_one

transfer a single **spi** transfer. - return 0 if the transfer is finished, - return 1 if the transfer is still in progress. When

the driver is finished with this transfer it must call `spi_finalize_current_transfer()` so the subsystem can issue the next transfer. Note: `transfer_one` and `transfer_one_message` are mutually exclusive; when both are set, the generic subsystem does not call your `transfer_one` callback.

handle_err

the subsystem calls the driver to handle an error that occurs in the generic implementation of `transfer_one_message()`.

mem_ops

optimized/dedicated operations for interactions with **SPI** memory. This field is optional and should only be implemented if the controller has native support for memory like operations.

cs_gpios

LEGACY: array of GPIO descs to use as chip select lines; one per CS number. Any individual value may be `-ENOENT` for CS lines that are not GPIOs (driven by the **SPI** controller itself). Use the `cs_gpiods` in new drivers.

cs_gpiods

Array of GPIO descs to use as chip select lines; one per CS number. Any individual value may be `NULL` for CS lines that are not GPIOs (driven by the **SPI** controller itself).

use_gpio_descriptors

Turns on the code in the **SPI** core to parse and grab GPIO descriptors rather than using global GPIO numbers grabbed by the driver. This will fill in `cs_gpiods` and `cs_gpios` should not be used, and SPI devices will have the `cs_gpiod` assigned rather than `cs_gpio`.

statistics

statistics for the **spi** controller

dma_tx

DMA transmit channel

dma_rx

DMA receive channel

dummy_rx

dummy receive buffer for full-duplex devices

dummy_tx

dummy transmit buffer for full-duplex devices

fw_translate_cs

If the boot firmware uses different numbering scheme what Linux expects, this optional hook can be used to translate between the two.

Description

Each **SPI** controller can communicate with one or more **spi_device** children. These make a small bus, sharing MOSI, MISO and SCK signals but not chip select signals. Each device may be configured to use a different clock rate, since those shared signals are ignored unless the chip is selected.

The driver for an **SPI** controller manages access to those devices through a queue of **spi_message** transactions, copying data between CPU memory and an SPI slave device. For each such message it queues, it calls the message's completion function when the transaction completes.

struct spi_res

spi resource management structure

Definition

```
struct spi_res {  
    struct list_head      entry;  
    spi_res_release_t release;  
    unsigned long long    data[];  
};
```

Members

entry

list entry

release

release code called prior to freeing this resource

data

extra data allocated for the specific use-case

Description

this is based on ideas from devres, but focused on life-cycle management during **spi_message** processing

struct spi_transfer

a read/write buffer pair

Definition

```
struct spi_transfer {  
    const void      *tx_buf;  
    void *rx_buf;  
    unsigned len;  
    dma_addr_t tx_dma;  
    dma_addr_t rx_dma;  
    struct sg_table tx_sg;  
    struct sg_table rx_sg;  
    unsigned cs_change:1;  
    unsigned tx_nbits:3;  
    unsigned rx_nbits:3;  
#define SPI_NBITS_SINGLE      0x01 ;  
#define SPI_NBITS_DUAL        0x02 ;  
#define SPI_NBITS_QUAD        0x04 ;  
    u8 bits_per_word;  
    u8 word_delay_usecs;  
    u16 delay_usecs;  
    u32 speed_hz;  
    u16 word_delay;  
    struct list_head transfer_list;  
};
```

Members

tx_buf

data to be written (dma-safe memory), or NULL

rx_buf

data to be read (dma-safe memory), or NULL

len

size of rx and tx buffers (in bytes)

tx_dma

DMA address of tx_buf, if **spi_message.is_dma_mapped**

rx_dma

DMA address of rx_buf, if **spi_message.is_dma_mapped**

tx_sg

Scatterlist for transmit, currently not for client use

rx_sg

Scatterlist for receive, currently not for client use

cs_change

affects chipselect after this transfer completes

tx_nbits

number of bits used for writing. If 0 the default (**SPI_NBITS_SINGLE**) is used.

rx_nbits

number of bits used for reading. If 0 the default (**SPI_NBITS_SINGLE**) is used.

bits_per_word

select a `bits_per_word` other than the device default for this transfer. If 0 the default (from `spi_device`) is used.

word_delay_usecs

microseconds to inter word delay after each word size (set by `bits_per_word`) transmission.

delay_usecs

microseconds to delay after this transfer before (optionally) changing the chipselect status, then starting the next transfer or completing this `spi_message`.

speed_hz

Select a speed other than the device default for this transfer. If 0 the default (from `spi_device`) is used.

word_delay

clock cycles to inter word delay after each word size (set by `bits_per_word`) transmission.

transfer_list

transfers are sequenced through `spi_message.transfers`

Description

SPI transfers always write the same number of bytes as they read. Protocol drivers should always provide `rx_buf` and/or `tx_buf`. In some cases, they may also want to provide DMA addresses for the data being transferred; that may reduce overhead, when the underlying driver uses dma.

If the transmit buffer is null, zeroes will be shifted out while filling `rx_buf`. If the receive buffer is null, the data shifted in will be discarded. Only “len” bytes shift out (or in). It’s an error to try to shift out a partial word. (For example, by shifting out three bytes with word size of sixteen or twenty bits; the former uses two bytes per word, the latter uses four bytes.)

In-memory data values are always in native CPU byte order, translated from the wire byte order (big-endian except with `SPI_LSB_FIRST`). So for example when `bits_per_word` is sixteen, buffers are 2N bytes long (`len = 2N`) and hold N sixteen bit words in CPU byte order.

When the word size of the **SPI** transfer is not a power-of-two multiple of eight bits, those in-memory words include extra bits. In-memory words are always seen by protocol drivers as right-justified, so the undefined (rx) or unused (tx) bits are always the most significant bits.

All **SPI** transfers start with the relevant chipselect active. Normally it stays selected until after the last transfer in a message. Drivers can affect the chipselect signal using `cs_change`.

- (i) If the transfer isn’t the last one in the message, this flag is used to make the chipselect briefly go inactive in the middle of the message. Toggling chipselect in this way may be needed to terminate a chip command, letting a single `spi_message` perform all of group of chip transactions together.
- (ii) When the transfer is the last one in the message, the chip may stay selected until the next transfer. On multi-device **SPI** busses with nothing blocking messages going to other devices, this is just a performance hint; starting a message to another device deselects this one. But in other cases, this can

be used to ensure correctness. Some devices need protocol transactions to be built from a series of `spi_message` submissions, where the content of one message is determined by the results of previous messages and where the whole transaction ends when the chipselect goes inactive.

When **SPI** can transfer in 1x, 2x or 4x. It can get this transfer information from device through **tx_nbits** and **rx_nbits**. In Bi-direction, these two should both be set. User can set transfer mode with `SPI_NBITS_SINGLE(1x)` `SPI_NBITS_DUAL(2x)` and `SPI_NBITS_QUAD(4x)` to support these three transfer.

The code that submits an **spi** message (and its `spi_transfers`) to the lower layers is responsible for managing its memory. Zero-initialize every field you don't set up explicitly, to insulate against future API updates. After you submit a message and its transfers, ignore them until its completion callback.

`struct spi_message`

one multi-segment **SPI** transaction

Definition

```
struct spi_message {
    struct list_head    transfers;
    struct spi_device   *spi;
    unsigned is_dma_mapped:1;
    void (*complete)(void *context);
    void *context;
    unsigned frame_length;
    unsigned actual_length;
    int status;
    struct list_head    queue;
    void *state;
    struct list_head    resources;
};
```

Members

`transfers`

list of transfer segments in this transaction

`spi`

SPI device to which the transaction is queued

`is_dma_mapped`

if true, the caller provided both dma and cpu virtual addresses for each transfer buffer

`complete`

called to report transaction completions

`context`

the argument to `complete()` when it's called

`frame_length`

the total number of bytes in the message

`actual_length`

the total number of bytes that were transferred in all successful segments

status

zero for success, else negative errno

queue

for use by whichever driver currently owns the message

state

for use by whichever driver currently owns the message

resources

for resource management when the **spi** message is processed

Description

A **spi_message** is used to execute an atomic sequence of data transfers, each represented by a struct **spi_transfer**. The sequence is “atomic” in the sense that no other spi_message may use that SPI bus until that sequence completes. On some systems, many such sequences can execute as a single programmed DMA transfer. On all systems, these messages are queued, and might complete after transactions to other devices. Messages sent to a given spi_device are always executed in FIFO order.

The code that submits an **spi_message** (and its spi_transfers) to the lower layers is responsible for managing its memory. Zero-initialize every field you don't set up explicitly, to insulate against future API updates. After you submit a message and its transfers, ignore them until its completion callback.

```
void spi_message_init_with_transfers(struct spi_message * m, struct spi_transfer * xfers, unsigned int num_xfers)
```

Initialize **spi_message** and append transfers

Parameters

```
struct spi_message * m
```

spi_message to be initialized

```
struct spi_transfer * xfers
```

An array of **spi** transfers

```
unsigned int num_xfers
```

Number of items in the xfer array

Description

This function initializes the given **spi_message** and adds each spi_transfer in the given array to the message.

```
struct spi_replaced_transfers
```

structure describing the **spi** transfer replacements that have occurred so that they can get reverted

Definition

```

struct spi_replaced_transfers {
    spi_replaced_release_t release;
    void *extradata;
    struct list_head replaced_transfers;
    struct list_head *replaced_after;
    size_t inserted;
    struct spi_transfer inserted_transfers[];
};

```

Members

release

some extra release code to get executed prior to relasing this structure

extradata

pointer to some extra data if requested or NULL

replaced_transfers

transfers that have been replaced and which need to get restored

replaced_after

the transfer after which the **replaced_transfers** are to get re-inserted

inserted

number of transfers inserted

inserted_transfers

array of **spi** transfers of array-size **inserted**, that have been replacing replaced_transfers

note

that **extradata** will point to **inserted_transfers**[**inserted]** if some extra allocation is requested, so alignment will be the same as for **spi** transfers

```

int spi_sync_transfer(struct spi_device * spi, struct spi_transfer * xfers, unsigned
int num_xfers)

```

synchronous **SPI** data transfer

Parameters

struct spi_device * spi

device with which data will be exchanged

struct spi_transfer * xfers

An array of **spi** transfers

unsigned int num_xfers

Number of items in the xfer array

Context

can sleep

Description

Does a synchronous **SPI** data transfer of the given spi_transfer array.

For more specific semantics see `spi_sync()`.

Return

Return: zero on success, else a negative error code.

```
int spi_write(struct spi_device * spi, const void * buf, size_t len)
```

SPI synchronous write

Parameters

```
struct spi_device * spi
```

device to which data will be written

```
const void * buf
```

data buffer

```
size_t len
```

data buffer size

Context

can sleep

Description

This function writes the buffer **buf**. Callable only from contexts that can sleep.

Return

zero on success, else a negative error code.

```
int spi_read(struct spi_device * spi, void * buf, size_t len)
```

SPI synchronous read

Parameters

```
struct spi_device * spi
```

device from which data will be read

```
void * buf
```

data buffer

```
size_t len
```

data buffer size

Context

can sleep

Description

This function reads the buffer **buf**. Callable only from contexts that can sleep.

Return

zero on success, else a negative error code.

```
ssize_t spi_w8r8(struct spi_device * spi, u8 cmd)
```

SPI synchronous 8 bit write followed by 8 bit read

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
u8 cmd
```

command to be written before data is read back

Context

can sleep

Description

Callable only from contexts that can sleep.

Return

the (unsigned) eight bit number returned by the device, or else a negative error code.

```
ssize_t spi_w8r16(struct spi_device * spi, u8 cmd)
```

SPI synchronous 8 bit write followed by 16 bit read

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
u8 cmd
```

command to be written before data is read back

Context

can sleep

Description

The number is returned in wire-order, which is at least sometimes big-endian.

Callable only from contexts that can sleep.

Return

the (unsigned) sixteen bit number returned by the device, or else a negative error code.

```
ssize_t spi_w8r16be(struct spi_device *spi, u8 cmd)
```

SPI synchronous 8 bit write followed by 16 bit big-endian read

Parameters

```
struct spi_device *spi
```

device with which data will be exchanged

```
u8 cmd
```

command to be written before data is read back

Context

can sleep

Description

This function is similar to **spi_w8r16**, with the exception that it will convert the read 16 bit data word from big-endian to native endianness.

Callable only from contexts that can sleep.

Return

the (unsigned) sixteen bit number returned by the device in cpu endianness, or else a negative error code.

```
struct spi_board_info
```

board-specific template for a **SPI** device

Definition

```

struct spi_board_info {
    char modalias[SPI_NAME_SIZE];
    const void      *platform_data;
    const struct property_entry *properties;
    void *controller_data;
    int irq;
    u32 max_speed_hz;
    u16 bus_num;
    u16 chip_select;
    u32 mode;
};

```

Members

modalias

Initializes `spi_device.modalias`; identifies the driver.

platform_data

Initializes `spi_device.platform_data`; the particular data stored there is driver-specific.

properties

Additional device properties for the device.

controller_data

Initializes `spi_device.controller_data`; some controllers need hints about hardware setup, e.g. for DMA.

irq

Initializes `spi_device.irq`; depends on how the board is wired.

max_speed_hz

Initializes `spi_device.max_speed_hz`; based on limits from the chip datasheet and board-specific signal quality issues.

bus_num

Identifies which `spi_controller` parents the `spi_device`; unused by `spi_new_device()`, and otherwise depends on board wiring.

chip_select

Initializes `spi_device.chip_select`; depends on how the board is wired.

mode

Initializes `spi_device.mode`; based on the chip datasheet, board wiring (some devices support both 3WIRE and standard modes), and possibly presence of an inverter in the chipselect path.

Description

When adding new `SPI` devices to the device tree, these structures serve as a partial device template. They hold information which can't always be determined by drivers. Information that `probe()` can establish (such as the default transfer wordsizes) is not included here.

These structures are used in two places. Their primary role is to be stored in tables of board-specific device descriptors, which are declared early in board initialization and then used (much later) to populate a controller's device tree after the that controller's driver initializes. A secondary (and atypical) role is as a parameter to `spi_new_device()` call, which happens after those controller drivers are active in some dynamic board configuration models.

```
int spi_register_board_info(struct spi_board_info const * info, unsigned n)
```

register **SPI** devices for a given board

Parameters

```
struct spi_board_info const * info
```

array of chip descriptors

```
unsigned n
```

how many descriptors are provided

Context

can sleep

Description

Board-specific early init code calls this (probably during arch_initcall) with segments of the **SPI** device table. Any device nodes are created later, after the relevant parent SPI controller (bus_num) is defined. We keep this table of devices forever, so that reloading a controller driver will not make Linux forget about these hard-wired devices.

Other code can also call this, e.g. a particular add-on board might provide **SPI** devices through its expansion connector, so code initializing that board would naturally declare its SPI devices.

The board info passed can safely be __initdata ... but be careful of any embedded pointers (platform_data, etc), they're copied as-is. Device properties are deep-copied though.

Return

zero on success, else a negative error code.

```
int __spi_register_driver(struct module * owner, struct spi_driver * drv)
```

register a **SPI** driver

Parameters

```
struct module * owner
```

owner module of the driver to register

```
struct spi_driver * drv
```

the driver to register

Context

can sleep

Return

zero on success, else a negative error code.

```
struct spi_device * spi_alloc_device(struct spi_controller * ctrl)
```

Allocate a new **SPI** device

Parameters

```
struct spi_controller * ctrl
```

Controller to which device is connected

Context

can sleep

Description

Allows a driver to allocate and initialize a **spi**_device without registering it immediately. This allows a driver to directly fill the **spi**_device with device parameters before calling

```
spi_add_device()
```

 on it.

Caller is responsible to call `spi_add_device()` on the returned **spi**_device structure to add it to the **SPI** controller. If the caller needs to discard the spi_device without adding it, then it should call

```
spi_dev_put()
```

 on it.

Return

a pointer to the new device, or NULL.

```
int spi_add_device(struct spi_device * spi)
```

Add **spi**_device allocated with spi_alloc_device

Parameters

```
struct spi_device * spi
```

spi_device to register

Description

Companion function to **spi**_alloc_device. Devices allocated with spi_alloc_device can be added onto the spi bus with this function.

Return

0 on success; negative errno on failure

```
struct spi_device * spi_new_device(struct spi_controller * ctrl, struct spi_board_info * chip)
```

instantiate one new **SPI** device

Parameters

```
struct spi_controller * ctrl
```

Controller to which device is connected

```
struct spi_board_info * chip
```

Describes the **SPI** device

Context

can sleep

Description

On typical mainboards, this is purely internal; and it's not needed after board init creates the hard-wired devices. Some development platforms may not be able to use `spi_register_board_info` though, and this is exported so that for example a USB or parport based adapter driver could add devices (which it would learn about out-of-band).

Return

the new device, or NULL.

```
void spi_unregister_device(struct spi_device * spi)
```

unregister a single **SPI** device

Parameters

```
struct spi_device * spi
```

`spi_device` to unregister

Description

Start making the passed **SPI** device vanish. Normally this would be handled by

```
spi_unregister_controller()
```

```
void spi_finalize_current_transfer(struct spi_controller * ctrl)
```

report completion of a transfer

Parameters

```
struct spi_controller * ctrl
```

the controller reporting completion

Description

Called by **SPI** drivers using the core `transfer_one_message()` implementation to notify it that the current interrupt driven transfer has finished and the next one may be scheduled.

```
struct spi_message * spi_get_next_queued_message(struct spi_controller * ctrl)
```

called by driver to check for queued messages

Parameters

```
struct spi_controller * ctrl
```

the controller to check for queued messages

Description

If there are more messages in the queue, the next message is returned from this call.

Return

the next message in the queue, else NULL if the queue is empty.

```
void spi_finalize_current_message(struct spi_controller * ctrl)
```

the current message is complete

Parameters

```
struct spi_controller * ctrl
```

the controller to return the message to

Description

Called by the driver to notify the core that the message in the front of the queue is complete and can be removed from the queue.

```
int spi_slave_abort(struct spi_device * spi)
```

abort the ongoing transfer request on an **SPI** slave controller

Parameters

```
struct spi_device * spi
```

device used for the current transfer


```
struct spi_controller * __spi_alloc_controller(struct device * dev, unsigned int size,
bool slave)
```

allocate an **SPI** master or slave controller

Parameters

struct device * dev

the controller, possibly using the platform_bus

unsigned int size

how much zeroed driver-private data to allocate; the pointer to this memory is in the driver_data field of the returned device, accessible with `spi_controller_get_devdata()`.

bool slave

flag indicating whether to allocate an **SPI** master (false) or SPI slave (true) controller

Context

can sleep

Description

This call is used only by **SPI** controller drivers, which are the only ones directly touching chip registers. It's how they allocate an `spi_controller` structure, prior to calling

`spi_register_controller()`.

This must be called from context that can sleep.

The caller is responsible for assigning the bus number and initializing the controller's methods before calling `spi_register_controller()`; and (after errors adding the device) calling `spi_controller_put()` to prevent a memory leak.

Return

the **SPI** controller structure on success, else NULL.

```
int spi_register_controller(struct spi_controller * ctrl)
```

register **SPI** master or slave controller

Parameters

struct spi_controller * ctrl

initialized master, originally from `spi_alloc_master()` or `spi_alloc_slave()`

Context

can sleep

Description

SPI controllers connect to their drivers using some non-**SPI** bus, such as the platform bus. The final stage of `probe()` in that code includes calling `spi_register_controller()` to hook up to this SPI bus glue.

SPI controllers use board specific (often SOC specific) bus numbers, and board-specific addressing for SPI devices combines those numbers with chip select numbers. Since SPI does not directly support dynamic device identification, boards need configuration tables telling which chip is at which address.

This must be called from context that can sleep. It returns zero on success, else a negative error code (dropping the controller's refcount). After a successful return, the caller is responsible for calling `spi_unregister_controller()`.

Return

zero on success, else a negative error code.

```
int devm_spi_register_controller(struct device *dev, struct spi_controller *ctlr)
```

register managed **SPI** master or slave controller

Parameters

```
struct device *dev
```

device managing **SPI** controller

```
struct spi_controller *ctlr
```

initialized controller, originally from `spi_alloc_master()` or `spi_alloc_slave()`

Context

can sleep

Description

Register a **SPI** device as with `spi_register_controller()` which will automatically be unregistered and freed.

Return

zero on success, else a negative error code.

```
void spi_unregister_controller(struct spi_controller *ctlr)
```

unregister **SPI** master or slave controller

Parameters

```
struct spi_controller * ctrl
```

the controller being unregistered

Context

can sleep

Description

This call is used only by **SPI** controller drivers, which are the only ones directly touching chip registers.

This must be called from context that can sleep.

Note that this function also drops a reference to the controller.

```
struct spi_controller * spi_busnum_to_master(u16 bus_num)
```

look up master associated with bus_num

Parameters

```
u16 bus_num
```

the master's bus number

Context

can sleep

Description

This call may be used with devices that are registered after arch init time. It returns a refcounted pointer to the relevant **spi**_controller (which the caller must release), or NULL if there is no such master registered.

Return

the **SPI** master structure on success, else NULL.

```
void * spi_res_alloc(struct spi_device * spi, spi_res_release_t release, size_t size, gfp_t gfp)
```

allocate a **spi** resource that is life-cycle managed during the processing of a spi_message while using spi_transfer_one

Parameters

```
struct spi_device * spi
```

the `spi` device for which we allocate memory

```
spi_res_release_t release
```

the release code to execute for this resource

```
size_t size
```

size to alloc and return

```
gfp_t gfp
```

GFP allocation flags

Return

the pointer to the allocated data

This may get enhanced in the future to allocate from a memory pool of the `spi_device` or `spi_controller` to avoid repeated allocations.

```
void spi_res_free(void * res)
```

free an `spi` resource

Parameters

```
void * res
```

pointer to the custom data of a resource

```
void spi_res_add(struct spi_message * message, void * res)
```

add a `spi_res` to the `spi_message`

Parameters

```
struct spi_message * message
```

the `spi` message

```
void * res
```

the `spi_resource`

```
void spi_res_release(struct spi_controller * ctrl, struct spi_message * message)
```

release all `spi` resources for this message

Parameters

```
struct spi_controller * ctrl
```

the `spi_controller`

```
struct spi_message * message
```

the `spi_message`

```
struct spi_replaced_transfers * spi_replace_transfers(struct spi_message * msg, struct
spi_transfer * xfer_first, size_t remove, size_t insert, spi_replaced_release_t release,
size_t extradatasize, gfp_t gfp)
```

replace transfers with several transfers and register change with `spi_message.resources`

Parameters

```
struct spi_message * msg
```

the `spi_message` we work upon

```
struct spi_transfer * xfer_first
```

the first `spi_transfer` we want to replace

```
size_t remove
```

number of transfers to remove

```
size_t insert
```

the number of transfers we want to insert instead

```
spi_replaced_release_t release
```

extra release code necessary in some circumstances

```
size_t extradatasize
```

extra data to allocate (with alignment guarantees of struct `spi_transfer`)

```
gfp_t gfp
```

gfp flags

Return

pointer to `spi_replaced_transfers`,
PTR_ERR(...) in case of errors.

```
int spi_split_transfers_maxsize(struct spi_controller * ctrl, struct spi_message * msg,
size_t maxsize, gfp_t gfp)
```

split `spi` transfers into multiple transfers when an individual transfer exceeds a certain size

Parameters

```
struct spi_controller * ctrl
```

the `spi_controller` for this transfer

```
struct spi_message * msg
```

the `spi_message` to transform

```
size_t maxsize
```

the maximum when to apply this

```
gfp_t gfp
```

GFP allocation flags

Return

status of transformation

```
int spi_setup(struct spi_device * spi )
```

setup **SPI** mode and clock rate

Parameters

```
struct spi_device * spi
```

the device whose settings are being modified

Context

can sleep, and no requests are queued to the device

Description

SPI protocol drivers may need to update the transfer mode if the device doesn't work with its default. They may likewise need to update clock rates or word sizes from initial values. This function changes those settings, and must be called from a context that can sleep. Except for `SPI_CS_HIGH`, which takes effect immediately, the changes take effect the next time the device is selected and data is transferred to or from it. When this function returns, the spi device is deselected.

Note that this call will fail if the protocol driver specifies an option that the underlying controller or its driver does not support. For example, not all hardware supports wire transfers using nine bit words, LSB-first wire encoding, or active-high chipselects.

Return

zero on success, else a negative error code.

```
void spi_set_cs_timing(struct spi_device * spi, u8 setup, u8 hold, u8 inactive_dly)
```

configure CS setup, hold, and inactive delays

Parameters

```
struct spi_device * spi
```

the device that requires specific CS timing configuration

```
u8 setup
```

CS setup time in terms of clock count

```
u8 hold
```

CS hold time in terms of clock count

```
u8 inactive_dly
```

CS inactive delay between transfers in terms of clock count

```
int spi_async(struct spi_device * spi, struct spi_message * message)
```

asynchronous **SPI** transfer

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
struct spi_message * message
```

describes the data transfers, including completion callback

Context

any (irqs may be blocked, etc)

Description

This call may be used in_irq and other contexts which can't sleep, as well as from task contexts which can sleep.

The completion callback is invoked in a context which can't sleep. Before that invocation, the value of message->status is undefined. When the callback is issued, message->status holds either zero (to indicate complete success) or a negative error code. After that callback returns, the driver which issued the transfer request may deallocate the associated memory; it's no longer in use by any SPI core or controller driver code.

Note that although all messages to a spi_device are handled in FIFO order, messages may go to different devices in other orders. Some device might be higher priority, or have various "hard" access time requirements, for example.

On detection of any fault during the transfer, processing of the entire message is aborted, and the device is deselected. Until returning from the associated message completion callback, no other spi_message queued to that device will be processed. (This rule applies equally to all the synchronous transfer calls, which are wrappers around this core asynchronous primitive.)

Return

zero on success, else a negative error code.

```
int spi_async_locked(struct spi_device * spi, struct spi_message * message)
```

version of spi_async with exclusive bus usage

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
struct spi_message * message
```

describes the data transfers, including completion callback

Context

any (irqs may be blocked, etc)

Description

This call may be used in_irq and other contexts which can't sleep, as well as from task contexts which can sleep.

The completion callback is invoked in a context which can't sleep. Before that invocation, the value of message->status is undefined. When the callback is issued, message->status holds either zero (to indicate complete success) or a negative error code. After that callback returns, the driver which issued the transfer request may deallocate the associated memory; it's no longer in use by any **SPI** core or controller driver code.

Note that although all messages to a **spi**_device are handled in FIFO order, messages may go to different devices in other orders. Some device might be higher priority, or have various "hard" access time requirements, for example.

On detection of any fault during the transfer, processing of the entire message is aborted, and the device is deselected. Until returning from the associated message completion callback, no other **spi**_message queued to that device will be processed. (This rule applies equally to all the synchronous transfer calls, which are wrappers around this core asynchronous primitive.)

Return

zero on success, else a negative error code.

```
int spi_sync(struct spi_device * spi, struct spi_message * message)
```

blocking/synchronous **SPI** data transfers

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
struct spi_message * message
```

describes the data transfers

Context

can sleep

Description

This call may only be used from a context that may sleep. The sleep is non-interruptible, and has no timeout. Low-overhead controller drivers may DMA directly into and out of the message buffers.

Note that the **SPI** device's chip select is active during the message, and then is normally disabled between messages. Drivers for some frequently-used devices may want to minimize costs of selecting a chip, by leaving it selected in anticipation that the next message will go to the same chip. (That may increase power usage.)

Also, the caller is guaranteeing that the memory associated with the message will not be freed before this call returns.

Return

zero on success, else a negative error code.

```
int spi_sync_locked(struct spi_device * spi, struct spi_message * message)
```

version of **spi_sync** with exclusive bus usage

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
struct spi_message * message
```

describes the data transfers

Context

can sleep

Description

This call may only be used from a context that may sleep. The sleep is non-interruptible, and has no timeout. Low-overhead controller drivers may DMA directly into and out of the message buffers.

This call should be used by drivers that require exclusive access to the **SPI** bus. It has to be preceded by a `spi_bus_lock` call. The SPI bus must be released by a `spi_bus_unlock` call when the exclusive access is over.

Return

zero on success, else a negative error code.

```
int spi_bus_lock(struct spi_controller * ctrl)
```

obtain a lock for exclusive **SPI** bus usage

Parameters

```
struct spi_controller * ctrl
```

SPI bus master that should be locked for exclusive bus access

Context

can sleep

Description

This call may only be used from a context that may sleep. The sleep is non-interruptible, and has no timeout.

This call should be used by drivers that require exclusive access to the **SPI** bus. The SPI bus must be released by a spi_bus_unlock call when the exclusive access is over. Data transfer must be done by spi_sync_locked and spi_async_locked calls when the SPI bus lock is held.

Return

always zero.

```
int spi_bus_unlock(struct spi_controller * ctrl)
```

release the lock for exclusive **SPI** bus usage

Parameters

```
struct spi_controller * ctrl
```

SPI bus master that was locked for exclusive bus access

Context

can sleep

Description

This call may only be used from a context that may sleep. The sleep is non-interruptible, and has no timeout.

This call releases an **SPI** bus lock previously obtained by an spi_bus_lock call.

Return

always zero.

```
int spi_write_then_read(struct spi_device * spi, const void * txbuf, unsigned n_tx, void * rxbuf, unsigned n_rx)
```

SPI synchronous write followed by read

Parameters

```
struct spi_device * spi
```

device with which data will be exchanged

```
const void * txbuf
```

data to be written (need not be dma-safe)

```
unsigned n_tx
```

size of txbuf, in bytes

```
void * rxbuf
```

buffer into which data will be read (need not be dma-safe)

```
unsigned n_rx
```

size of rxbuf, in bytes

Context

can sleep

Description

This performs a half duplex MicroWire style transaction with the device, sending txbuf and then reading rxbuf. The return value is zero for success, else a negative errno status code. This call may only be used from a context that may sleep.

Parameters to this routine are always copied using a small buffer; portable code should never use this for more than 32 bytes. Performance-sensitive or bulk transfer code should instead use **spi**_{*async*,*sync*}() calls with dma-safe buffers.

Return

zero on success, else a negative error code.